

METRICS, EXPERIMENTATION & GTM

career coach.ai · PM Capstone Document · v1.0

Document Title	career coach.ai — Metrics, Experimentation & GTM
Product	career coach.ai
Owner (PM)	Abhinav Kotar
Version	v1.0
Status	In Review
Created On	24-05-2026
Last Updated	24-05-2026

1. North Star Metric

North Star Metric

Weekly Active Users completing 3 or more sessions

Unique users who complete at least 3 full coaching sessions within a rolling 7-day window.

1.1 Why This Metric Matters

Session completion is a meaningful proxy for real coaching value. a user who merely logs in has not experienced the product. Three sessions is the threshold at which the memory store has enough signal to meaningfully personalise subsequent coaching, which is the core differentiator of career coach.ai. A user who reaches this threshold has experienced the product at its best. Tracking this weekly rather than monthly catches engagement drop-off early enough to act on it.

1.2 User Value view

Three completed sessions represent a meaningful commitment to preparation. At this point the user has received structured feedback, seen their competency scores update, and experienced at least one session personalised to their prior performance. This is the moment users first understand why career coach.ai is different from a generic AI tool — memory makes the third session noticeably better than the first.

1.3 Business Value view

Weekly coached users drive every downstream business outcome: word-of-mouth referrals to batchmates, institutional interest from bootcamps, and eventual conversion to a paid tier. A product with high weekly coaching engagement is demonstrably working which is the primary signal needed to justify continued investment, expand to new cohorts, and build the institutional sales motion.

2. Supporting KPIs

The following seven KPIs sit beneath the North Star Metric, each measuring a distinct dimension of product health. Together they form an early warning system — a drop in any one KPI points to a specific part of the product that needs attention.

Metric	What It Measures	Why It Matters
Session Completion Rate	% of started sessions that reach the session summary screen.	The coaching loop only delivers value when completed. High abandonment mid-session signals friction in the answer-evaluation flow or response latency issues.
Competency Score Improvement Rate	Average delta in skill scores between a user's 1st and 5th session, per competency dimension.	The product's core promise is improvement. If scores are not moving after 5 sessions, the coaching content or memory personalisation is not working.
Day-7 Retention Rate	% of new users who complete at least one session in the 7 days following their first session.	Early signal of habit formation. A healthy D7 retention rate confirms the product created enough value in the first session to pull users back.
Resume Review Adoption Rate	% of active users who use the resume review feature at least once within their first 30 days.	Measures breadth of engagement beyond mock interviews. Low adoption signals either low feature awareness or low perceived value of the resume module.
Feedback Helpfulness Rate	% of session feedback cards rated 'Helpful' by users across all sessions.	Direct user signal on AI output quality. The most actionable quality metric available without human review of every session.
Session-to-Session Memory Utilisation	% of sessions where the vector store successfully loaded prior user context vs. cold-start fallback.	Measures the technical health of the core differentiator. If memory is failing silently, users are not experiencing the personalisation the product promises.
Monthly Active Users (MAU)	Unique users who complete at least one full coaching session in a calendar month.	Top-of-funnel growth signal. Tracks whether the user base is expanding month-over-month and provides the denominator for retention and engagement rate calculations.

3. AARRR Funnel Mapping

The AARRR framework maps each stage of the user lifecycle to a specific user action and a measurable metric, giving the product team a shared language for diagnosing where users are dropping off.

Funnel Stage	User Action	Metric
Acquisition	User discovers career coach.ai via a bootcamp referral, LinkedIn post, or organic search and signs up for an account.	Sign-up conversion rate (% of landing page visitors who create an account). Target: 25%+ in first cohort.
Activation	User completes their first full mock interview session and views their session summary with competency scores.	First-session completion rate (% of sign-ups who complete session 1 within 48 hours). Target: 60%+.
Retention	User returns and completes a second and third session within 7 days, experiencing memory-driven personalisation for the first time.	Day-7 retention rate (% of activated users who complete 3+ sessions within 7 days). Target: 40%+ at MVP stage.

Funnel Stage	User Action	Metric
Revenue	User converts from free tier to paid subscription, or an institution purchases a cohort licence.	Free-to-paid conversion rate. Not applicable in v1 — product is free during MVP validation phase.
Referral	User shares career coach.ai with a batchmate or colleague via word-of-mouth, direct link share, or social post.	Referral rate (% of active users who share a referral link within their first 30 days). Target: 20%+ in bootcamp cohorts.

Revenue note: The Revenue stage is intentionally inactive in v1. The MVP validation phase runs on a fully free model to maximise adoption. Monetisation is introduced in v2 once the core coaching loop is validated and retention benchmarks are met.

4. A/B Test Design

The following test targets the highest-impact drop-off point: users who complete session 1 but do not return for session 2. Addressing this gap is the single most leveraged intervention for improving the North Star Metric.

Hypothesis	Showing users a personalised 'Next session preview' card at the end of session 1 — displaying the specific competency the AI has identified as their weakest and the question type that will be focused on next — will increase Day-2 return rate compared to a generic 'Start your next session' CTA.
Control (A)	Session summary screen ends with a generic CTA button: 'Start your next session →'. No preview of what the next session will focus on. This is the current baseline behaviour.
Variant (B)	Session summary screen ends with a personalised preview card: 'Your next session will focus on Product Sense — your lowest-scoring area. We'll start with estimation questions at medium difficulty.' Followed by CTA: 'Start focused session →'.
Target Users	All new users who complete their first full mock interview session during the test period. Minimum sample size: 200 users per variant (400 total) to achieve 80% statistical power at 95% confidence interval, assuming a baseline D2 return rate of 35%.
Primary Metric	Day-2 return rate — % of users in each variant who complete at least one session within 48 hours of completing session 1.
Secondary Metrics	Session 2 completion rate, session 3 completion rate (does the effect persist), and session 2 helpfulness rating (does personalised framing raise or lower satisfaction with the coaching received).
Success Criteria	Variant B produces a statistically significant improvement in Day-2 return rate of 5 percentage points or more (e.g. 35% to 40%+) at $p < 0.05$, with no degradation in session 2 helpfulness rating vs. control.

5. Dashboard Plan

The dashboard plan defines what is measured, who needs to see it, and how often — ensuring the right people are looking at the right signals at the right cadence.

Dashboard Section	Metrics Shown	User of Dashboard	Frequency of Review
North Star Tracker	Weekly active coached users (3+ sessions); week-over-week trend; cohort breakdown by source.	PM, Founder	Daily (launch) / Weekly (steady state)
Funnel Health	Sign-up to activation to D7 retention rates; drop-off % at each stage; session 1 completion rate.	PM, Growth	Weekly
AI Quality Monitor	Feedback helpfulness rate; low-confidence badge trigger rate; schema validation failure rate; memory utilisation rate.	PM, Engineering	Daily
User Progress Metrics	Average competency score improvement (session 1 to session 5); % of users showing positive delta; skill radar distribution across user base.	PM, Product	Weekly
Feature Adoption	Resume review adoption rate; daily coaching loop usage; session type breakdown (behavioural vs. product sense vs. estimation).	PM, Product	Weekly
Growth & Referral	MAU trend; referral link shares; referral conversion rate; new user sources (direct, referral, organic, partner).	PM, Growth	Weekly
Error & Reliability	LLM API error rate; P99 response latency; memory retrieval timeout rate; session abandonment due to errors.	Engineering	Real-time alerts + Daily review

6. GTM Strategy

6.1 Target Audience

Primary: Learners enrolled in outcome-focused bootcamps and accelerated programmes in India (Masai, upGrad, Scaler, Newton School) — aged 20–27, actively job-seeking within 3–6 months. These users have high preparation urgency, peer networks that accelerate word-of-mouth, and programme coordinators who can champion the tool at the cohort level.

Secondary: Self-directed career switchers (engineering to PM, ops to PM) preparing for interviews outside of a formal programme. Reached via LinkedIn, PM communities (PMDojo, Product School), and organic search.

6.2 Positioning Statement

For aspiring product managers and early-career professionals who cannot afford a human coach, career coach.ai is the only interview preparation tool that remembers your progress and gets smarter about your specific gaps with every session — so your preparation compounds instead of resetting.

6.3 Launch Channels

- **Bootcamp Partnerships (Primary):** Direct outreach to programme coordinators at Masai, upGrad, and Scaler. Offer the tool free to one cohort in exchange for user feedback and a case study. Programme coordinators become internal champions who drive sign-up urgency among learners.
- **LinkedIn Content (Organic):** Weekly posts targeting PM job seekers: mock interview teardowns, competency framework explainers, and 'what good looks like' examples. Builds trust and drives direct sign-ups without paid spend.
- **PM Community Seeding:** Share the product in Product School India, PMDojo, and relevant LinkedIn/WhatsApp groups with a free access link. Target active members who are currently interviewing — highest urgency, most authentic early reviews.
- **Referral Programme:** Built-in referral link in the product from day one. Users who refer 3 batchmates unlock an additional feature (company-specific question pack preview). Peer referral is the most trusted channel in a bootcamp cohort context.
- **SEO / Organic Search:** Long-tail content targeting 'PM interview prep', 'CIRCLES framework practice', 'product sense interview questions India'. Low-cost, compounding channel — not a launch driver but valuable within 3–6 months.

6.4 Key Message

'The more you use it, the better it knows you.' Every other prep tool forgets you the moment you close the tab. career coach.ai remembers every session, every weak spot, and every improvement — so your coach gets sharper every time you show up.

6.5 Launch Checklist

- Core product functional: mock interview, memory store, session summary, and skill radar all working end-to-end.
- AI quality baseline confirmed: feedback helpfulness rate >70% on internal test sessions before external users are onboarded.
- Privacy policy and AI disclosure copy finalised and live on the landing page.
- Onboarding flow tested with 5 internal users — profile creation to first session summary under 10 minutes.
- Analytics instrumentation verified: all AARRR funnel events firing correctly in PostHog/Mixpanel.
- Bootcamp partner confirmed: at least one cohort coordinator committed to sharing the product with their learners.
- Referral mechanism live: referral link generation and tracking functional before first external user is onboarded.
- Error alerting configured: LLM API failures, memory retrieval timeouts, and schema validation errors all trigger Slack alerts.
- Feedback collection live: per-session helpfulness rating and 'Report this feedback' path functional.
- Launch comms ready: LinkedIn post, partner email template, and community seeding message drafted and reviewed.

7. Pricing and Monetisation

7.1 Recommended Model: Freemium with Institutional Licensing

career coach.ai will launch on a **fully free model** during the MVP validation phase. No payment wall, no session limits. The goal in v1 is adoption, retention data, and proof of coaching effectiveness — not revenue. Introducing pricing before the core loop is validated would suppress the adoption needed to generate the data required to improve it.

Tier	Price	Features	Target
Free (v1 MVP)	Unlimited	All MVP features: mock interviews, memory store, resume review, skill radar, session summaries.	All users during MVP validation. No expiry.
Pro (v2)	~Rs. 499/month (~\$6)	Everything in Free + company-specific question packs, LinkedIn profile review, priority support, unlimited session history.	Active users who have completed 5+ sessions and are actively interviewing.
Institutional (v2)	Custom (per cohort)	Cohort dashboard for coordinators, bulk onboarding, aggregated progress reports, dedicated onboarding session.	Bootcamp coordinators and programme managers at partner institutions.

7.2 Why This Model Fits the Product

- **Free removes the adoption barrier** that would otherwise prevent the data flywheel from starting. Memory-based personalisation requires session data — which only accumulates if users actually use the product.
- **Pro pricing is affordable for the target user** — Rs. 499/month is less than a single hour with a human coach on Topmate, and directly competitive with the user's stated willingness to pay.
- **Institutional licensing is the highest-leverage monetisation channel** — one cohort contract covers 50–200 users, is recurring, and comes with a built-in adoption mechanism through the programme coordinator.

8. Post-Launch Feedback Plan

Feedback collection is designed to be continuous, multi-channel, and directly actionable. Every feedback mechanism maps to a specific product decision — no feedback is collected without a clear owner and a defined path to action.

8.1 Feedback Channels

Channel	What Is Collected
In-product (passive)	Per-session helpfulness rating (Helpful / Not helpful + optional reason). Session-level star rating (1–5). 'Report this feedback' flag on every AI response. Automatic low-confidence badge triggers logged as implicit feedback.
In-product (active)	Post-session survey triggered after session 3: 'Is the coaching feeling more personalised than your first session?' Weekly prompt to power users (5+ sessions): 'What is the one thing that would make this more useful?'
User interviews	Structured 30-minute interviews with 5–8 users at the end of the first 4 weeks. Focus on: activation experience, perceived coaching quality, memory personalisation awareness, and willingness to refer.
Community & social	Monitor mentions in PM communities and WhatsApp cohort groups. Track organic shares and referral link usage patterns. Direct outreach to users who share without a referral link.
Support tickets	All inbound support messages reviewed weekly. Categorised by theme: AI quality, technical issue, feature request, onboarding confusion. Themes surfaced in weekly product review.

8.2 How Feedback Influences Future Iterations

Feedback Signal	Product Response
Helpfulness rate drops below 70% on any question type	Immediate prompt engineering review for that question type. New prompt variant tested within 5 business days.
Post-session 3 survey: >30% say coaching does not feel personalised	Memory retrieval logic reviewed. Consider surfacing 'Based on your last session...' language more explicitly in session opening.
User interview insight: users unaware of memory feature	Onboarding flow updated with explicit memory explanation. In-session message added: 'I remember you struggled with X last time — let us work on that.'
Support theme: confusion about competency scores	Tooltip copy on skill radar improved. FAQ entry added. Potential onboarding walkthrough of the dashboard added to v2 backlog.
Referral rate below 15% after 4 weeks	Referral UX reviewed. Referral prompt timing tested (post-session 3 vs. post-session 1). Referral incentive strengthened.
MAU growth flat after week 4	Re-engagement email sequence designed for users inactive for 48+ hours. Session streak mechanic added to product backlog for v2.