

UX DESIGN — USER FLOW & WIREFRAMES

career coach.ai · PM Capstone Document · v1.0

Document Title	career coach.ai — UX Design: User Flow & Wireframes
Product	career coach.ai
Owner (PM)	Abhinav Kotar
Version	v1.0
Status	In Review
Created On	25-05-2026
Last Updated	24-05-2026

1. User Flow


The user flow maps the complete journey a learner takes through career coach.ai — from first landing on the product to completing a coached session and returning for the next one. The flow covers both the mock interview path and the resume review path, which share the same entry point and home dashboard but diverge after that. Edge cases (memory retrieval failure, LLM timeout, harmful content detection) are shown as branching paths that resolve gracefully without breaking the main journey.

1.1 Main Journey Overview

Stage	Steps
Entry	Landing Page → Sign Up / Log In
Onboarding	Role selection → Experience level → AI Disclosure Modal (must acknowledge before proceeding)
Home Dashboard	Skill radar + session streak + daily coaching prompt + resume review entry point
Path A	Start Mock Interview → Session type selection → Memory loaded → Question generated → User answers → Guardrail check → AI evaluates → Feedback card shown → Loop (3–5 questions) → Session Summary → Rate session → Memory updated → Back to Dashboard
Path B	Upload Resume → File parsed → AI generates section feedback → Feedback panel shown → Back to Dashboard
Edge Cases	Memory retrieval failure → silent fallback to DB profile. LLM timeout (>10s) → fallback question bank. Harmful content detected → blocked + logged. User flags feedback → escalation queue.

1.2 User Flow Diagram

The diagram below shows the complete user flow including decision points, edge case branches, and the memory update loop. Refer to the full annotated diagram for node-level detail.




2. Wireframes

Three low-fidelity wireframes are provided below, covering the three most critical screens in the product. Each screen represents a distinct moment in the user journey — habit formation, core value delivery, and retention trigger — and together they communicate the full coaching loop without requiring any other screens to be understood.

2.1 Wireframe 1 — Home Dashboard


The home dashboard is the screen users return to every day. It surfaces three things at a glance: where the user currently stands (skill radar), whether they are maintaining their habit (streak), and what they should do

next (daily coaching prompt). The memory system is most visible here — the radar updates after every session and the daily prompt is generated from the user's lowest competency score, making it immediately clear the product knows them.




2.2 Wireframe 2 — Live Session Screen

The live session screen is the core value delivery moment. The question is shown prominently at the top with its competency dimension labelled so the user knows what is being assessed. The answer input is large and uncluttered — no distractions during practice. Once submitted, the feedback card appears in the same view with dimension scores, a qualitative tip, and a 'what good looks like' benchmark. The AI label and confidence badge (if triggered) are shown persistently on the feedback card, and the helpful/not helpful rating is directly below it



2.3 Wireframe 3 — Session Summary Screen

The session summary is the retention-driving moment. It shows the user exactly how this session compared to their last one — trend arrows (up/down/flat) next to each competency score make improvement immediately visible. Three ranked action items tell the user precisely what to work on before the next session. The 'next session preview' card at the bottom (showing the competency that will be focused on next) creates a pull to return. The session star rating is the final interaction before returning to the dashboard.



3. UX Explanation

The design decisions in career coach.ai are not aesthetic choices — each one is a direct response to a specific user behaviour or trust problem identified during research. This section explains the reasoning behind the flow, the feature placement, and the trust and safety elements built into the interface.

3.1 Why the Flow Is Designed This Way

The flow is built around a single principle: **reduce the distance between opening the app and doing something useful**. Every screen before the coaching session is a necessary gate — not extra steps. The AI disclosure modal on onboarding is mandatory because trust must be established before the user encounters their first AI-generated evaluation. The dashboard is the first screen after login (not a welcome page or tutorial) because returning users should reach their progress and their next action in one tap.

The two paths — mock interview and resume review — are separated at the dashboard level rather than nested inside each other. This is deliberate: a user who comes back for a quick 20-minute session should not have to navigate past resume features to get there. Each path is self-contained and returns to the same dashboard, keeping the mental model simple.

3.2 How the Design Makes the Product Easy to Use

Design Feature	Usability Rationale
Daily coaching prompt on dashboard	Eliminates decision fatigue. The user does not need to choose what to practise — the AI has already identified the highest-priority focus area and surfaced it as the primary CTA. One tap starts the right session.
Question and feedback on the same screen	Users do not lose context switching between screens. The question stays visible above the feedback card so the user can re-read it while processing the evaluation. No pagination, no tabs.
Trend arrows on session summary	Progress is made concrete and immediate. Instead of requiring users to compare numbers across sessions themselves, the interface does the work — a green up arrow on 'Structure' tells the story instantly.
Next session preview card	Removes the barrier to returning. By telling the user exactly what comes next before they leave, the product creates a natural pull rather than relying on the user to self-motivate a return visit.
Streak counter	Low-effort habit reinforcement. The streak is visible on the dashboard and increments immediately after session completion. It rewards consistency without requiring any additional user action.

3.3 How the Design Builds Trust in AI Output

Trust in AI output is not built through a single disclaimer — it is built through a set of consistent, layered signals that appear throughout the interface. The following elements work together to ensure users engage with the AI's feedback critically rather than blindly.

Trust Element	How It Works
AI label on every feedback card	A persistent 'AI-generated feedback' label in the top corner of every evaluation response ensures users never forget the source of the coaching. It is non-dismissible and appears regardless of feedback quality.
Low-confidence badge	When the model's response contains more than 3 hedging phrases, a yellow badge appears: 'Lower confidence — use as a starting point.' This teaches users to calibrate their trust based on the specific output, not apply blanket trust or distrust to all AI feedback.
Benchmark example (collapsed by default)	Every feedback card includes a 'What good looks like' example answer for the question. Showing the standard the AI is evaluating against makes the scoring feel transparent and earned rather than arbitrary.
Helpful / Not helpful rating	Placing the rating directly below the feedback card (not on a separate survey screen) signals that the user's opinion of the AI's output matters and will influence future sessions. This is a trust-building mechanism as much as a feedback mechanism.
Trend arrows vs. absolute scores	Showing directional change (up/down/flat) alongside absolute scores reduces the risk of users over-indexing on a single session score. The interface implicitly tells users to trust the trend, not any one number.

3.4 Privacy, Safety, and Feedback Elements in the Interface

Privacy and safety are surfaced at the moments of highest risk — not buried in a settings page. The following elements are visible in the interface at the points where user data or AI behaviour could create concern.

Interface Element	Category	What It Does
AI Disclosure Modal (onboarding)	Privacy / Trust	Shown on first login, must be acknowledged before the user can proceed. Explains that all coaching is AI-generated, how session data is stored, and that resume files are processed in-session only by default. Sets expectations before any data is shared.
Resume upload consent notice	Privacy	A banner above the resume upload panel confirms: 'Your resume will be processed this session only. It will not be stored unless you choose to save it.' Explicit, in-context, one-tap to proceed.
'Report this feedback' link	Safety / Feedback	Available on every feedback card and in the session summary footer. Opens a short form: 'What was wrong?' Escalates to human review queue. Users are told their report will be reviewed — not sent into a void.
Helpful / Not helpful rating	Feedback	Persistent below every AI evaluation. One-tap interaction. 'Not helpful' opens an optional reason selector (Too generic / Incorrect / Irrelevant / Other). Logged and used for prompt quality monitoring.
Session star rating	Feedback	Shown at the end of every session before returning to the dashboard. 1–5 stars. Sessions rated 1–2 stars are automatically flagged for PM review. Takes 2 seconds — low friction, high signal.
Low-confidence badge	Safety / Trust	Automatically appended to any response where the model hedges excessively. Tells the user to treat the output cautiously without requiring them to interpret the language themselves.
Data deletion note in settings	Privacy	Account settings surface a clear statement: 'Deleting your account will permanently remove all session data and memory within 30 days.' No small print, no buried toggles.

All privacy, safety, and feedback elements are designed to be visible at the moment of relevance — not surfaced only in a settings page or terms document. The principle is: if a user encounters a moment where their data or trust is at stake, the relevant control or disclosure is already on that screen.