

AI TOOL USAGE APPENDIX

career coach.ai · PM Capstone Document · v1.0

Document Title	career coach.ai — AI Tool Usage Appendix
Product	career coach.ai
Owner (PM)	Abhinav Kotar
Version	v1.0
Status	In Review
Created On	29-05-2026
Last Updated	29-05-2026

1. Purpose of This Document

This appendix provides a transparent record of every AI tool used in the creation of the career coach.ai PM Capstone. It documents what each tool was used for, what was asked of it, and — critically how the output was reviewed, challenged, or modified before being incorporated into the final deliverables. The goal is to demonstrate that AI tools were used as thinking aids and productivity accelerators, not as a replacement for product judgment.

All strategic decisions in this project selection, MVP scoping, metric choices, pricing model, GTM channel prioritisation, and UX flow structure were made by the PM. AI tools generated drafts, surfaced options, and accelerated documentation. I have directed, challenged, and approved every substantive output.

2. Tools Used — Summary

Tool	Primary Purpose	Phase
Groq	Real-time trending project research	Project selection phase
Perplexity	Resume-value and job-market project shortlisting	Project selection phase
Claude	All document drafting, structuring, and content generation across 5 documents	Throughout all phases
Canva	User flow diagram creation	UX design phase
Whimsical	Low-fidelity wireframe creation	UX design phase

3. Detailed AI Tool Usage

The table below documents each distinct use of an AI tool across the capstone. Rows with dashed borders are reserved for you to complete with tools or tasks not yet recorded.

Tool Used	Purpose	Prompt Used	How You Reviewed or Modified the Output
Groq	Real-time trending AI project research to identify which projects are gaining traction in the market right now.	"What are the most trending AI-powered products and project ideas in 2025 that would be valuable on a PM resume? Focus on real usage growth and job market relevance."	Reviewed the output against personal interest and feasibility. Cross-referenced with Perplexity findings. Used as a longlist only — final shortlist was a deliberate personal decision based on resume fit, not Groq's ranking.
Perplexity	Evaluating which shortlisted projects would have the most practical resume value and relevance for PM roles specifically.	"Which of these AI project ideas would be most impressive on a product management resume in India: RAG-based tutor, AI career coach, medical scribe, drug discovery tool, or UI automation agent? Consider PM interview relevance, differentiation, and buildability."	Output was used as one input alongside a direct conversation with Claude about project trade-offs. Final project decision (AI career coach with memory) was made after explicitly pushing back on the initial recommendation and stress-testing the differentiation angle.
Claude	Evaluating and shortlisting capstone project options. Stress-testing differentiation of the AI career coach vs. existing RAG tutor on resume.	"I already have a RAG agent AI tutor on my resume. Go through the shortlist again and find the next best project that adds something new."	Challenged Claude's first recommendation (Agent TARS integration) as scope creep. Pushed back and received a revised recommendation to focus on adaptive memory as the differentiating feature rather than browser automation.
Claude	Drafting 01_Product_Discovery_and_Strategy — all 10 sections including personas,	"Create a document called 01_Product_Discovery_and_Strategy.pdf. Include: Product Idea Overview,	Reviewed all 10 sections. Requested

Tool Used	Purpose	Prompt Used	How You Reviewed or Modified the Output
Claude	JTBD, value proposition, MoSCoW table.	Problem Discovery, Target User Segment, JTBD Statement, User Personas (x2), User Research Insights, Product Vision, Value Proposition, MVP Scope, Feature Prioritisation using MoSCoW."	reformatting to match Campus Connect PRD style (white background, blue headings, structured tables). Persona details and pain points were accepted as realistic. MoSCoW prioritisation reviewed and approved without changes.
	Drafting 02_PRD_and_Agile_Backlog — product objective, user stories, acceptance criteria, functional and non-functional requirements, assumptions, dependencies, risks, agile backlog.	"Create 02_PRD_and_Agile_Backlog. Include: Product Objective, Problem and Goal, User Stories (3–5)(relatable to a student, a grad and a career switcher), Acceptance Criteria (Given/When/Then), Functional Requirements, Non-Functional Requirements, Assumptions, Dependencies, Risks, Agile Backlog with 1 Epic and story points."	Reviewed all sections. Accepted Given/When/Then format for acceptance criteria. Agile backlog story points (13 for memory store, 8 for mock interview) reviewed and accepted as reasonable. Document format matched Campus Connect style as requested.
	Drafting 03_AI_Feature_and_Responsible_AI — AI capability mapping, input/output table, 9-step workflow, guardrails, responsible AI risks, transparency framework.	"Write 03_AI_Feature_and_Responsible_AI. Use model-agnostic LLM references (Claude / GPT-4o). Guardrails section should be specific — describe actual validation logic and thresholds. Treat resume data at capstone level (lighter). Include sections 8 (Guardrails) and 9 (AI Transparency) as additions after initial draft."	Three pre-build clarification questions were asked and answered before drafting began. Sections 8 and 9 were added as a follow-up after reviewing the initial 7-section draft. Guardrail logic (confidence threshold >3 hedging phrases, schema validation, 0.75 cosine similarity threshold) was reviewed and accepted as technically credible.
Claude	Proposing and validating KPIs for 04_Metrics_Experimentation_and_GTM before drafting began.	"Check in with me first — what 6 KPIs are you going to use for the supporting metrics?" [Claude proposed 6 KPIs; PM	Reviewed the proposed KPI list. Rejected

Tool Used	Purpose	Prompt Used	How You Reviewed or Modified the Output
		reviewed and responded: "maybe add MAU, other than that looks good go forward."]	DAU/MAU ratio and time-on-platform as vanity metrics. Added MAU as a 7th KPI with a specific definition (users completing at least one full session per calendar month, not just logins). This was a direct PM decision that changed the document.
Claude	Drafting 04_Metrics_Experimentation_and_GTM — North Star Metric, KPI table, AARRR funnel, A/B test design, dashboard plan, GTM strategy, pricing model, post-launch feedback plan.	"Go forward with the doc — include North Star Metric(crosscheck with different metric and conclude), Supporting KPIs (7 including MAU), Create an AARRR Funnel, A/B Test Design, Dashboard Plan, GTM Strategy With a proper positioning statement and launch channels(include PM communities), Pricing/Monetisation, Post-Launch Feedback Plan. Word format."	Reviewed all 8 sections. Pricing model (free MVP → freemium v2 → institutional licensing) was accepted as aligned with the product's equity-access mission. A/B test design (session 1 end screen, Day-2 return rate as primary metric) reviewed and accepted. GTM channels reviewed — bootcamp partnerships confirmed as primary channel.
	Generating the complete user flow text structure for use in Canva, and proposing the three wireframe screens to build.	"Give me clean text for the user flow — I'll draw it in Figma/Whimsical. Render wireframes inline in chat first so I can review before deciding on format."	Reviewed the proposed 3-screen wireframe selection (Dashboard, Live Session, Session Summary). Confirmed these were the right screens. Reviewed the inline Dashboard wireframe rendered in chat. Decided to use own wireframes in Whimsical rather than the inline version. Used the

Tool Used	Purpose	Prompt Used	How You Reviewed or Modified the Output
Claude			text user flow chart from Claude as the basis for the Canva diagram.
	Drafting Section 3 (UX Explanation) of 05_UX_and_User_Flow — design rationale, usability decisions, trust-building elements, privacy and safety interface elements.	"Refer to the screens generically — explain how we put different features and why. Build into Word doc with 4 diagram placeholders: 1 full-width for user flow + 3 individual wireframe placeholders."	Reviewed UX explanation sections. The privacy/safety table with colour-coded categories (indigo for privacy, yellow for safety, green for feedback) was accepted as clearly structured. Placeholder formatting (dashed blue borders with labels) confirmed as correct before build.
Canva	Creating the visual user flow diagram for 05_UX_and_User_Flow.	Used the structured text user flow generated by Claude in this conversation as the source. Recreated the flow as a visual diagram in Canva using the step-by-step node structure provided.	Visual layout, node labels, and branching logic reviewed against the original text flow for accuracy. Edge cases (memory failure, LLM timeout, harmful content detection) confirmed as included in the diagram.
Whimsical	Creating the three low-fidelity wireframes (Dashboard, Live Session, Session Summary) for 05_UX_and_User_Flow.	Used the wireframe descriptions and feature rationale from Claude's inline Dashboard wireframe and the UX Explanation section (Sections 2.1, 2.2, 2.3) as the brief for building in Whimsical.	Wireframes reviewed against the feature rationale in Section 3 of the UX document. Confirmed that AI label, confidence badge, helpful/not helpful rating, trend arrows, and next session preview card were all represented in the appropriate screens.

4. Reflection : PM Decisions vs. AI-Generated Output

A transparent AI usage log requires more than listing what was generated — it requires an honest account of where human judgment shaped the outcome and where AI was accepted without significant challenge. The following table draws that distinction explicitly.

Decision / Output	Origin	Detail
Project selection: AI career coach with adaptive memory	PM Decision	Multiple tools (Groq, Perplexity, Claude) were consulted. The final choice was made after explicitly rejecting Claude's first alternative recommendation (Agent TARS integration) as scope creep and requesting a reframe around a single differentiating feature.
Dropping TARS integration in favour of memory as the differentiator	PM Decision	Claude initially recommended combining TARS UI automation with the career coach. The PM pushed back, calling it 'scope creep dressed up as a feature.' Claude revised the recommendation. This was a direct product judgment call.
Adding MAU as a 7th KPI	PM Decision	Claude proposed 6 KPIs and asked for review before proceeding. The PM added MAU with a specific definition (session-completing users, not logins). Small change, but a deliberate one that shaped the metrics framework.
Rejecting DAU/MAU ratio and time-on-platform as vanity metrics	PM Decision	Claude's initial KPI reasoning was reviewed and two commonly used metrics were explicitly rejected as not meaningful for this product. This reflects product judgment about what actually matters.
Requesting Campus Connect document format	PM Decision	The visual format, heading style, table structure, and overall document aesthetic were directed by the PM with reference to an existing document. Claude implemented the format but did not choose it.
Persona names, pain points, and research insights	AI-Generated (vetted & accepted)	Personas (Riya Sharma, Arjun Mehta) and research insights were generated by Claude and accepted without significant modification. They were reviewed for realism and relevance but not rewritten.
Guardrail logic (confidence threshold, schema validation)	AI-Generated (vetted & accepted)	The specific technical guardrail design — >3 hedging phrases threshold, 0.75 cosine similarity, JSON-mode output enforcement — was proposed by Claude and accepted as technically credible. Not independently verified.
GTM channel prioritisation (bootcamp partnerships as primary)	AI-Generated (vetted & accepted)	Claude proposed bootcamp partnerships as the primary GTM channel based on the target user profile. This was accepted as logical given the Masai context but was not independently validated with market research.
A/B test design (Day-2 return rate, session 1 end screen)	AI-Generated (vetted & accepted)	The A/B test hypothesis, variant design, and success criteria were proposed by Claude and accepted with minor review. Sample size

Decision / Output	Origin	Detail
		calculation (200 per variant) was accepted without independent statistical verification.
Wireframe screen selection (Dashboard, Live Session, Summary)	PM Decision	Claude proposed these three screens and explained the rationale. The PM confirmed the selection as correct — this was an active agreement, not passive acceptance.

Origin key: Green = PM Decision (actively directed or changed) · Yellow = AI-Generated and accepted (reviewed but not materially modified)