



CareerCoach.ai

Levelling the Playing Field with Agentic AI

Abhinav Kotar | Product Manager

The Macro Funnel

India's Tech Talent Pool:
1.5M+ Annual Engineering Grads

Non-IIT
Placement Rate:
<40%

The gap is rarely knowledge; it is
preparation quality and practice access.

The Micro Inequity

Human
Coach

\$30-\$100/hr,
scheduled, scarce

Generic
Generative
AI

Free, but forgets you
instantly, repeats
generic advice, zero
progress tracking

Riya Sharma, 23 | Aspiring PM

Final-year Tier-2 engineering student.

Pain Point:

Repeats the same weak frameworks without improving; receives no actionable feedback.

Current Workaround:

Passive YouTube videos, generic ChatGPT prompts with no session continuity.

Arjun Mehta, 26 | Career Switcher

IT Developer switching to Product.

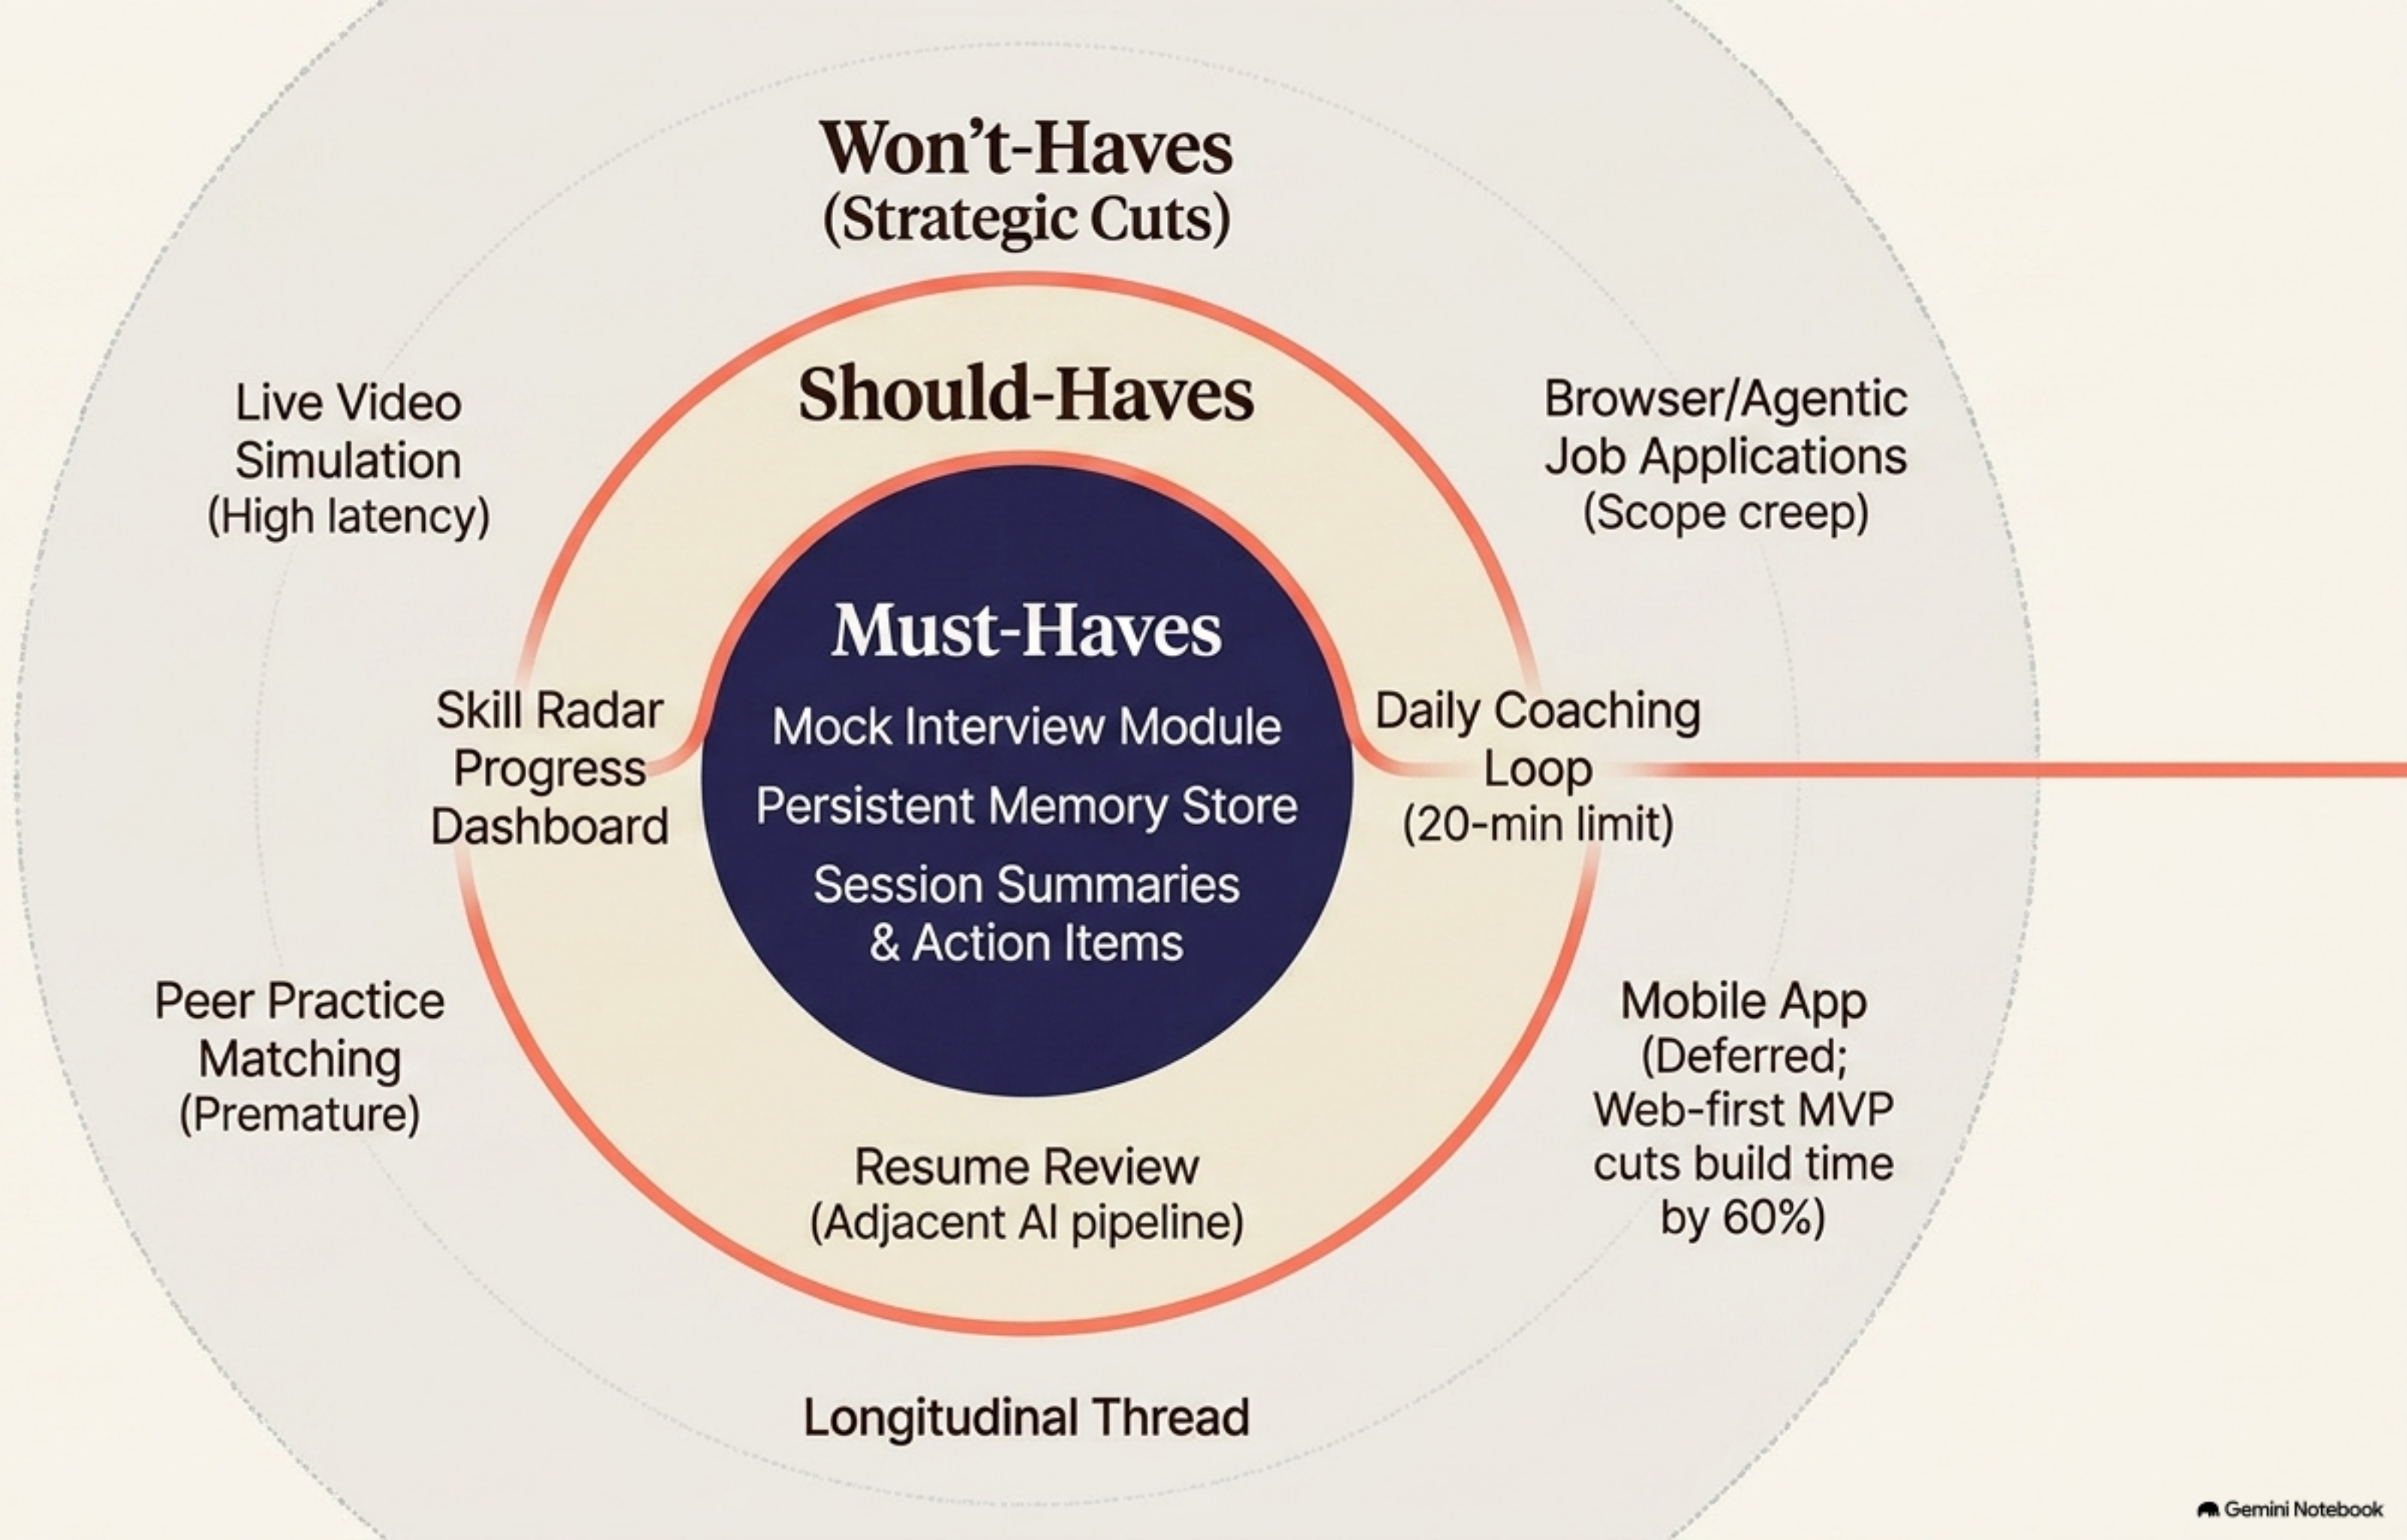
Pain Point:

Struggles to translate engineering experience into PM narratives; severely time-constrained (1 hr/day).

Current Workaround:

Sporadic practice, uncompleted online courses.

When I am practicing for a high-stakes interview, I want an adaptive AI to leverage my session history to drill into my specific weak areas, so I can systematically close my preparation gaps.



Latency Mitigation:
Pre-fetching

User Login

Context
fetched before
session starts

Vector Database



Pinecone /
ChromaDB
stores session
embeddings

Context Retrieval



Cosine Similarity > 0.75
Ensures only relevant
mistakes are pulled

Context Injection

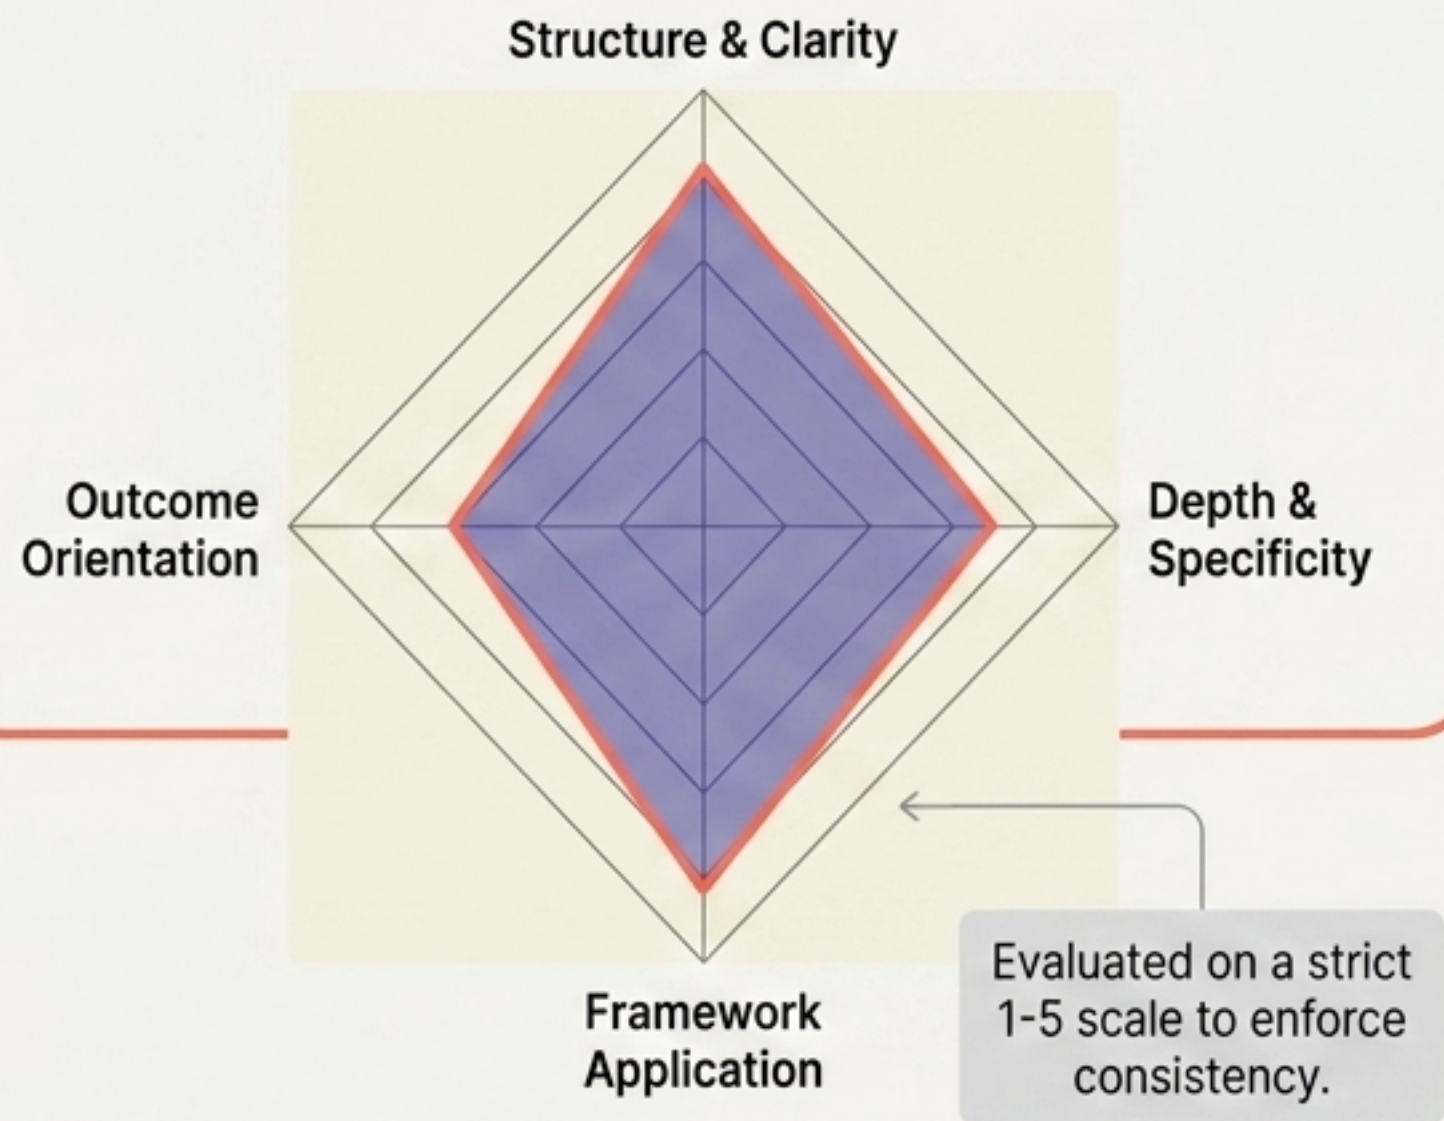
Prior weak areas
& action items
packaged into
prompt

LLM Prompt Execution

Dynamically
adapts question
difficulty based
on injected memory

Dynamic Evaluation & Guardrails

The 4-D Evaluation Rubric



The Guardrail Scanner

Format Guardrail:

JSON-Mode Enforced
(Rejects malformed outputs)

Confidence Scanner:

Hedging Check Active. Scans for phrases like 'I think', 'possibly'.
>3 instances trigger visual warning.

⚠ Low-Confidence Badge

Intercepts output, warning user to treat feedback as directional.

Trust-Centered UX: Reducing Friction & Building Habit

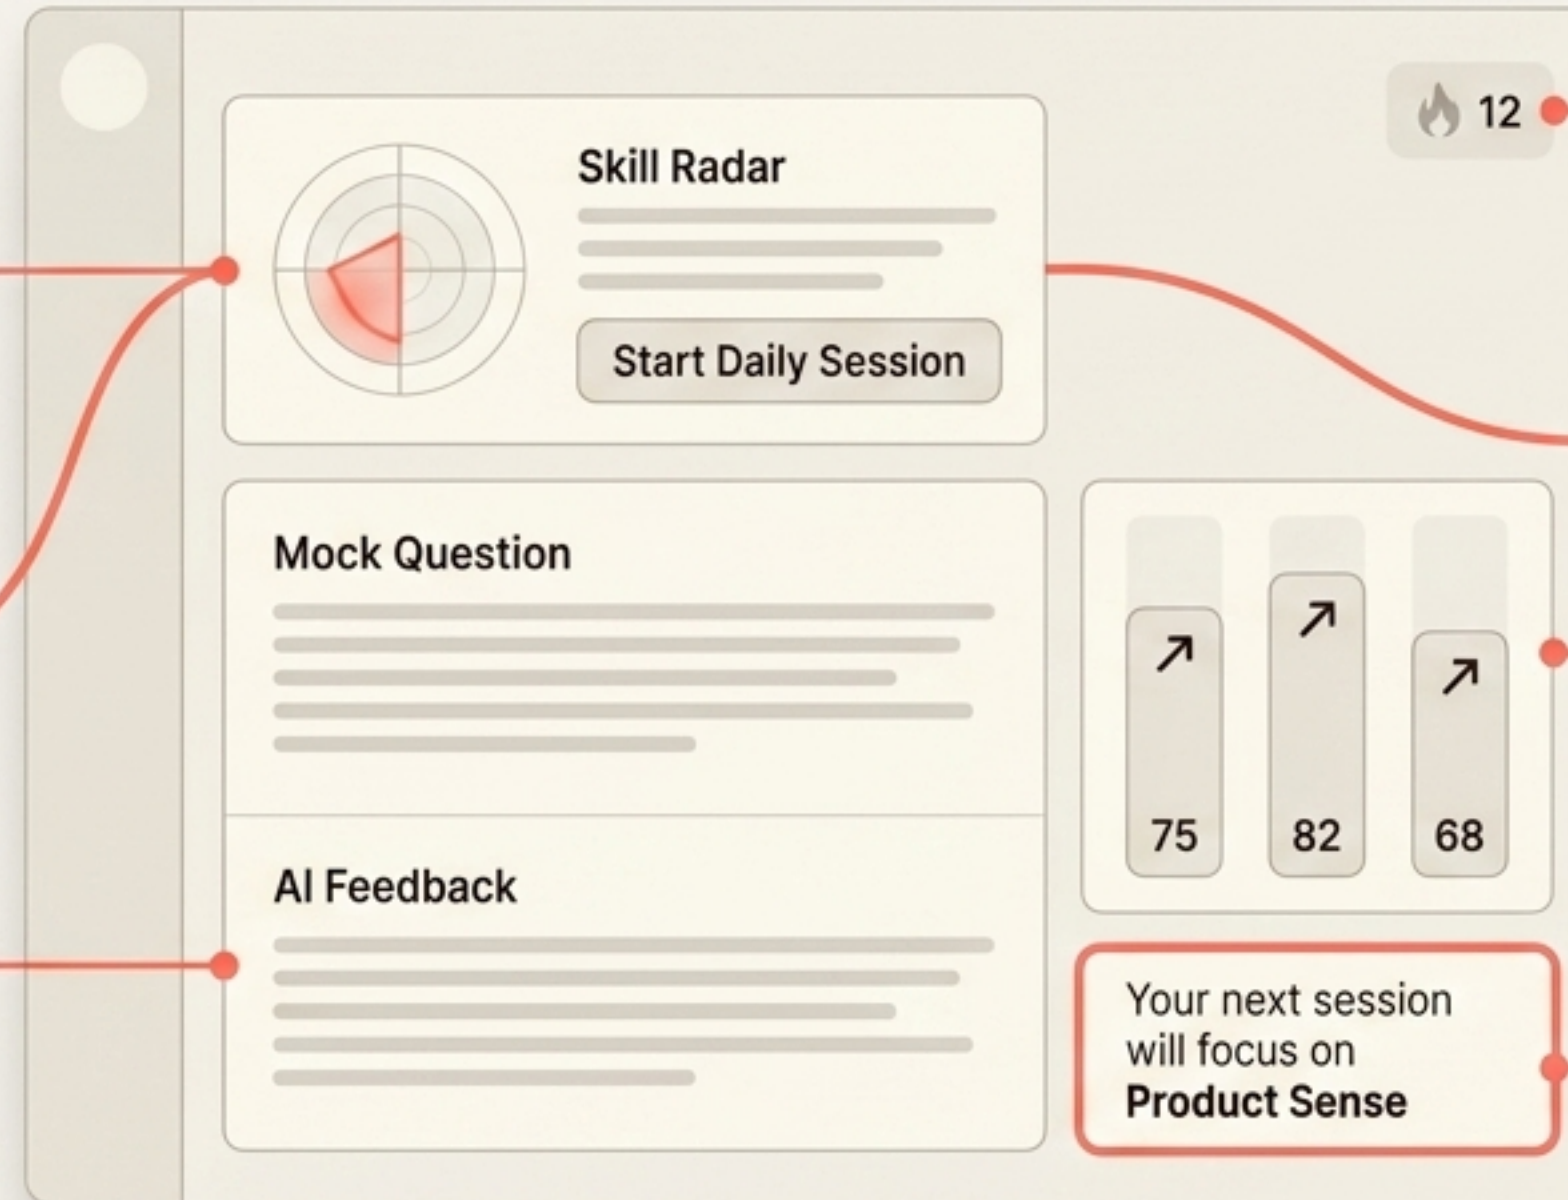
Eliminates decision fatigue.

The AI identifies the lowest-scoring competency and auto-generates the session.

The Longitudinal Thread

Question and feedback share the same screen.

No pagination, preserving mental context.



Streak Counter:

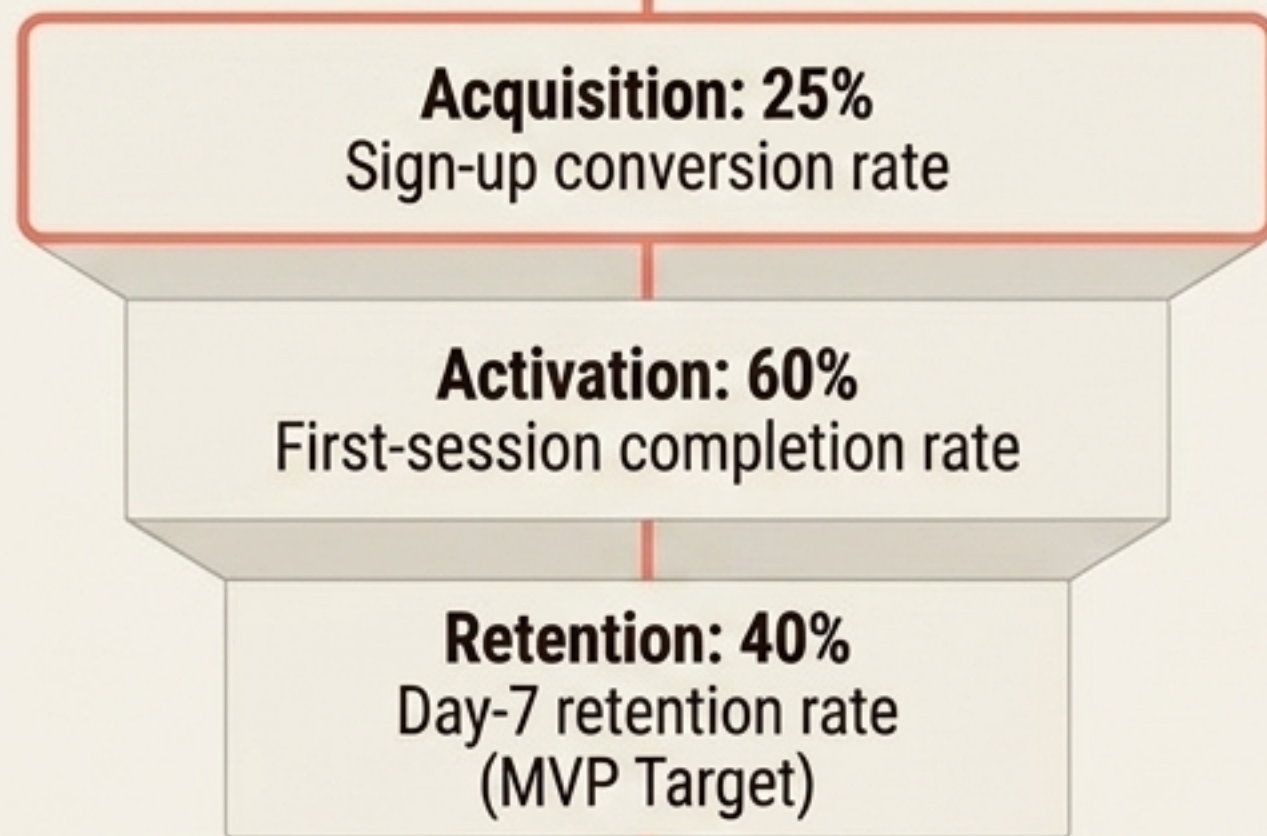
Low-friction habit reinforcement.

Abstracts raw scores into

directional momentum to make progress tangible.

Behavioral trigger.

Tells the user exactly what they will conquer tomorrow before they close the app today.



The North Star Weekly Active Users Completing 3+ Sessions

Why 3 sessions? This is the exact threshold where longitudinal memory kicks in. Session 3 is noticeably smarter than Session 1. Logging in is vanity; completing 3 sessions is value.

The Metrics Graveyard

~~DAU/MAU Ratio~~

~~Raw Time on Platform~~

Rejected: Metrics that reward friction instead of coaching efficiency.

Targeting the Day-1 Drop-Off Leak

Control A (Baseline)

Start your next session →

Variant B (The Intervention)



Your next session will focus on **Product Sense** – your lowest-scoring area.

Target Metric

Day-2 Return Rate

Sample Size

n=400

(80% statistical power
at 95% confidence)

Success Threshold

+5% increase

(without degrading session
helpfulness scores)

Stage 1: Free/MVP Validation

Fully free, unlimited sessions.
Organic LinkedIn acquisition.

Purpose: Eliminate adoption barriers to rapidly acquire session data, building the memory flywheel.

Stage 2: Institutional Partnerships

B2B2C distribution. Target Logos: Masai, Scaler, Newton School.

Purpose: Tap into high-urgency, outcome-focused bootcamps as a primary growth engine.

Stage 3: Premium Monetization

Rs. 499/month Pro Tier.

Purpose: Anchor pricing aggressively below a single human coaching session (\$30+).

Responsible AI Architecture: The Shield.

Input (Privacy & Data Minimization)

Session-only resume processing.

Resumes are processed dynamically in-session and wiped by default. Explicit opt-in required for persistent storage.

System (Bias & Fairness)

Decoupling language fluency from substance.

Rubric explicitly prevents penalizing non-native non-native English speakers. "Clarity" is strictly defined as logical flow and structure, not grammatical perfection.

Output (Recourse & Human Loop)

The Human Escalation Queue

"Report this feedback" buttons on every UI card. Sessions rated 1-2 stars bypass AI and immediately escalate to a PM review queue to correct prompt drift.

The Default Path

Product Strategy

Build a generic ChatGPT wrapper.

Feature Scoping

Say “yes” to scope creep (e.g., integrating Agent TARS for browser automation).

Defining Success

Optimize for vanity (DAU/MAU, total logins).

The PM Leadership Path

Product Strategy

Build a longitudinal memory engine that compounds user value over time.

Feature Scoping

Ruthlessly reject adjacent features to protect the core coaching loop and MVP timeline.

Defining Success

Optimize for actualized value (WAUs completing 3+ personalized sessions).

We don't just manage AI products. We actively steer them to deliver equitable, compounding human value.